

ALI RAZA

alirazabhutta.10@gmail.com - aliraza@bu.edu - <https://razaaliraza.github.io/>

OVERVIEW

Interests: Operating Systems, Computer Networks, Cloud Computing, Distributed Systems, Performance Analysis and Tuning
Skills: Linux Kernel Programming, eBPF, Systems Programming, C, Python, Bash

EDUCATION

Sep 2017 - Mar 2023	Ph.D. Computer Science BOSTON UNIVERSITY Advisor: Prof. Orran Krieger
Sep 2014 - Jun 2016	M.S. Computer Science LAHORE UNIVERSITY OF MANAGEMENT SCIENCES, PAKISTAN Advisor: Prof. Ihsan Ayyub Qazi
Sep 2008 - Jun 2012	B.S. Electrical Engineering LAHORE UNIVERSITY OF MANAGEMENT SCIENCES, PAKISTAN Senior Year Project Supervisor: Prof. Sohaib Khan

EXPERIENCE

Mar 2023 - Present	Senior Engineer, HPC System Architecture SAMSUNG SEMICONDUCTOR INC., USA Responsibilities: <ul style="list-style-type: none">• Research on optimizing operating systems to improve performance for traditional HPC and new AI/ML workloads.• Research on novel supercomputer architectures to identify opportunities to co-optimize hardware and software.• Publish the research as scholarly articles and file patents.
--------------------	--

SELECTED PROJECTS

PhD Thesis - Unikernel Linux (UKL)

This research, which I lead, explores if unikernel techniques (e.g., highly optimized transitions between kernel and application code, customized code paths, run-to-completion, link-time optimization of kernel and application code, etc.) can be integrated into Linux while preserving its battle-tested code, development community, and ecosystem of tools, applications, and supported hardware. Our prototype, called Unikernel Linux (UKL) demonstrates that it is possible, and can support a large class of unmodified applications and hardware, and result in significant performance improvements. Complex unmodified applications (e.g., Redis, Memcached) can improve by over 10% in 99th tail latency and throughput, and small modifications to the applications can result in more than 20% improvement. A group of Red Hat Linux developers is working with us to submit the core UKL 550 line patch upstream; the full set of optimizations is just over 1200 LoC. We are exploring a broad set of use cases ranging from accelerat-

ing cloud applications and core infrastructure services, to enabling new device driver models.

Github Link: <https://github.com/unikernelLinux/ukl>

Elastic Secure Datacenter

This research, which I participated in at the start of my PhD, explored how we can securely and rapidly multiplex physical servers between many different tenants while minimizing trust in the provider. The Bolted prototype we developed enabled tenants to use existing unmodified provisioning systems for HPC, enterprise and cloud services. Security sensitive tenants can deploy their own attestation services to ensure that previous tenants did not compromise the firmware. Keylime attestation service used in Bolted has developed a large open source community and is being used in both Red Hat products and IBM's public cloud. The fundamental mechanisms from this work have now been adapted by the OpenStack community, and are currently being integrated and tested at the MGHPC data center.

Context Aware WiFi Bitrate Adaptation

Indoor wireless channel conditions suffer from signal attenuation, multi-path interference and shadowing effects, decreasing Wi-Fi throughput, especially when high data rates are used. On the contrary, in outdoor settings, better channel conditions allow higher bit rates, thus improving throughput. For mobile devices, especially those which frequently transition between indoor and outdoor settings, WiFi bitrates need to quickly adapt to channel conditions. But the WiFi protocol (IEEE 802.11 family) estimates channel conditions over timescales of hundreds of packets or needs to be trained to understand the SNR-BER relationship, and is better suited to stationary devices. As part of my MS thesis, I designed an indoor-outdoor detection system which uses sensors present in mobile devices such as accelerometer, light sensor and GPS etc., and use this information to quickly adapt WiFi bitrate. The results show faster WiFi bitrate adaptation and better average throughput for a mobile device.

PUBLICATIONS

Unikernel Linux (UKL) - EuroSys '23 [PDF]

This work explores if unikernel techniques can be integrated into a general-purpose OS while preserving its battle-tested code, development community, and ecosystem of tools, applications, and hardware support. Our prototype demonstrates both a path to integrate unikernel techniques in Linux and that such techniques can result in performance advantages for unmodified applications. Expert developers can modify the application to call internal kernel functionality and optimize across the application/kernel boundary for more significant gains. The changes to the Linux kernel are modest (1250 LOC).

Unikernels: The Next Stage of Linux's Dominance - HotOS '19 [PDF]

Unikernels have demonstrated advantages over Linux in many important domains, but have lagged in adoption and ease of use. They lack the battle-tested code base and large open source community which Linux has. In this paper, we explore if Linux can be built into a unikernel, and would that give us any performance benefits. We build an early Linux unikernel prototype and demonstrate that some simple changes can bring dramatic performance advantages.

Using SGX-Based Virtual Clones for IoT Security - IEEE NCA '18 [PDF]

Widespread permeation of IoT devices into our daily lives has created a diverse spectrum of security and privacy concerns unique to the IoT ecosystem. Conventional host and network security mechanisms fail to address these issues due to resource

constraints, ad-hoc network models and vendor-centric data collection and sharing policies. Hence, there is a need to redesign the IoT infrastructure to secure both the device and the data. To this end, we propose a design where users are in the driving seat, devices are less exposed and data sharing models are flexible and fine-grained. Our proposal comprises hardware-secured data banks based on Intel Software Guard Extensions (SGX) to house the data in clouds without the need to trust the cloud provider. Virtual clones (shadows) of devices running on top of these data banks serve as competent proxies of actual IoT devices hiding away device weaknesses. The proposed infrastructure is scalable and robust and serves as a good first step for the community to build on and improve.

A Secure Cloud with Minimal Provider Trust - HotCloud 18 [\[PDF\]](#)

In this paper we propose Bolted which is a new architecture for a bare metal cloud with the goal of providing security-sensitive customers of a cloud the same level of security and control that they can obtain in their own private data centers. It allows tenants, rather than the provider, to control the trade-offs between security, price, and performance. A prototype demonstrates scalable end-to-end security with a small overhead compared to a less secure alternative.

It's All in the Name: Why Some URLs are More Vulnerable to Typosquatting - IEEE INFOCOM 2018 [\[PDF\]](#)

Typosquatting is a blackhat practice that relies on human error and low-cost domain registrations to hijack legitimate traffic from well-established websites. We explore the relationship between human hand anatomy, keyboard layouts and typing mistakes to understand why some URLs are more vulnerable to typing mistakes than others. Furthermore, we predict the most likely typos for each URL which can then be defensively registered.

An Anomaly Detection Fabric for Clouds Based on Collaborative VM Communities - CCGrid 2017 [\[PDF\]](#)

To mitigate security threats in the cloud, we propose a hypervisor layer anomaly detection system based on system call monitoring, which compresses the stream of system calls at their source making the system scalable and near real-time. Not requiring modifications to the guest OS or the application make it ideal for the datacenter setting. Additionally, for robust and early detection of threats, we share information about attacks in their early stages to provide immunity to the entire deployment.

TALKS AND PRESENTATIONS

- DevConf.us 2021 [video](#)
- DevConf.us 2020 [video](#)
- Red Hat Research Day 2020 [video](#)
- DevConf.us 2019 [video](#)
- DevConf.us 2018 [video](#)